

MLPH Final Project - Lung Cancer Prediction Models

Farah Beche and Han Oo

2024-05-13

```
#Import Dataset
raw_data <- read.csv("/Users/farahbeche/Downloads/lungcancerdataset.csv")
# Explore dataset
head(raw_data)
```

```
##  GENDER AGE SMOKING YELLOW_FINGERS ANXIETY PEER_PRESSURE CHRONIC.DISEASE
## 1      M  69      1              2      2              1              1
## 2      M  74      2              1      1              1              2
## 3      F  59      1              1      1              2              1
## 4      M  63      2              2      2              1              1
## 5      F  63      1              2      1              1              1
## 6      F  75      1              2      1              1              2
##  FATIGUE ALLERGY WHEEZING ALCOHOL.CONSUMING COUGHING SHORTNESS.OF.BREATH
## 1      2      1      2              2      2              2
## 2      2      2      1              1      1              2
## 3      2      1      2              1      2              2
## 4      1      1      1              2      1              1
## 5      1      1      2              1      2              2
## 6      2      2      2              1      2              2
##  SWALLOWING.DIFFICULTY CHEST.PAIN LUNG_CANCER
## 1      2      2      YES
## 2      2      2      YES
## 3      1      2      NO
## 4      2      2      NO
## 5      1      1      NO
## 6      1      1      YES
```

```
summary(raw_data)
```

```
##      GENDER      AGE      SMOKING      YELLOW_FINGERS
## Length:309      Min.   :21.00      Min.   :1.000      Min.   :1.00
## Class :character 1st Qu.:57.00      1st Qu.:1.000      1st Qu.:1.00
## Mode  :character Median :62.00      Median :2.000      Median :2.00
##              Mean  :62.67      Mean  :1.563      Mean  :1.57
##              3rd Qu.:69.00      3rd Qu.:2.000      3rd Qu.:2.00
##              Max.   :87.00      Max.   :2.000      Max.   :2.00
##      ANXIETY      PEER_PRESSURE      CHRONIC.DISEASE      FATIGUE
## Min.   :1.000      Min.   :1.000      Min.   :1.000      Min.   :1.000
## 1st Qu.:1.000      1st Qu.:1.000      1st Qu.:1.000      1st Qu.:1.000
## Median :1.000      Median :2.000      Median :2.000      Median :2.000
## Mean   :1.498      Mean   :1.502      Mean   :1.505      Mean   :1.673
```

```
## 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:2.000
## Max. :2.000 Max. :2.000 Max. :2.000 Max. :2.000
## ALLERGY WHEEZING ALCOHOL.CONSUMING COUGHING
## Min. :1.000 Min. :1.000 Min. :1.000 Min. :1.000
## 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:1.000
## Median :2.000 Median :2.000 Median :2.000 Median :2.000
## Mean :1.557 Mean :1.557 Mean :1.557 Mean :1.579
## 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:2.000
## Max. :2.000 Max. :2.000 Max. :2.000 Max. :2.000
## SHORTNESS.OF.BREATH SWALLOWING.DIFFICULTY CHEST.PAIN LUNG_CANCER
## Min. :1.000 Min. :1.000 Min. :1.000 Length:309
## 1st Qu.:1.000 1st Qu.:1.000 1st Qu.:1.000 Class :character
## Median :2.000 Median :1.000 Median :2.000 Mode :character
## Mean :1.641 Mean :1.469 Mean :1.557
## 3rd Qu.:2.000 3rd Qu.:2.000 3rd Qu.:2.000
## Max. :2.000 Max. :2.000 Max. :2.000
```

Observations:

- Predictors are all binary data, except for the patient's age.
- Response variable is 'LUNG_CANCER'
- The binary predictors use different ways to represent the data
 - LUNG_CANCER is represented by YES, NO
 - GENDER is represented by M(male), F(female)
 - SMOKING, YELLOW_FINGERS, ANXIETY, etc. are represented by 1(NO), 2(YES)

```
# Explore dataset
# Check for missing values
missing_values <- colSums(is.na(raw_data))
missing_values
```

```
## GENDER AGE SMOKING
## 0 0 0
## YELLOW_FINGERS ANXIETY PEER_PRESSURE
## 0 0 0
## CHRONIC.DISEASE FATIGUE ALLERGY
## 0 0 0
## WHEEZING ALCOHOL.CONSUMING COUGHING
## 0 0 0
## SHORTNESS.OF.BREATH SWALLOWING.DIFFICULTY CHEST.PAIN
## 0 0 0
## LUNG_CANCER
## 0
```

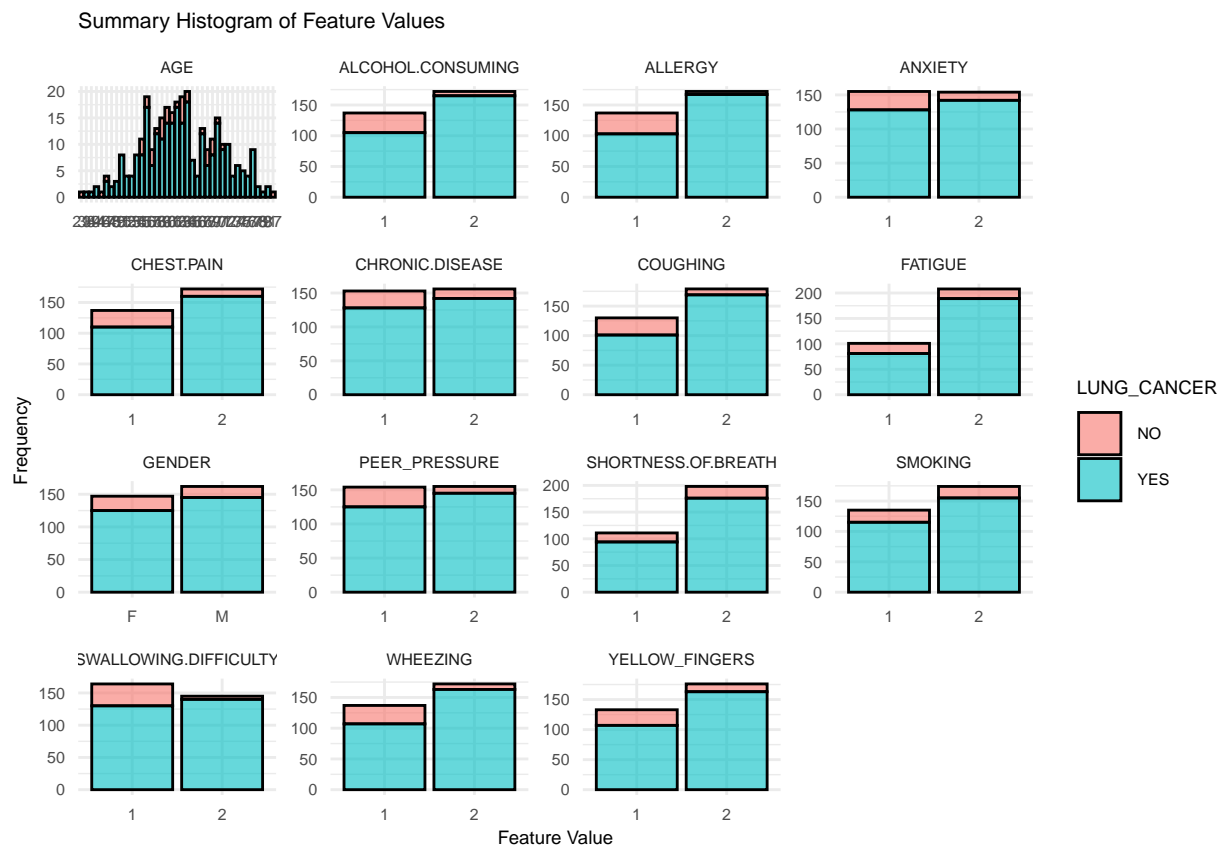
- There are no missing values in any of the variables.

```

# Explore dataset
# Summary histogram of feature values for each predictor variable
summary_histogram <- raw_data %>%
  gather(key = "Variable", value = "Value", -LUNG_CANCER) %>%
  ggplot(aes(x = Value, fill = as.factor(LUNG_CANCER))) +
  geom_bar(stat = "count", color = "black", alpha = 0.6) +
  facet_wrap(~ Variable, scales = "free") +
  labs(title = "Summary Histogram of Feature Values",
       x = "Feature Value",
       y = "Frequency",
       fill = "LUNG_CANCER") +
  theme_minimal() +
  theme(text = element_text(size = 7))

print(summary_histogram)

```



```

#Data Cleaning
# Encoding GENDER, assigning 0 to "M" and 1 to "F"
# Encoding LUNG_CANCER, assigning 0 to "NO" and 1 to "YES"
# Create a new dataset 'data' with encoded variables
raw_data$GENDER <- as.numeric(ifelse(raw_data$GENDER == "M", 0, 1))
raw_data$LUNG_CANCER <- as.numeric(ifelse(raw_data$LUNG_CANCER == "NO", 0, 1))

data <- raw_data
head(data)

```

```

##      GENDER AGE SMOKING YELLOW_FINGERS ANXIETY PEER_PRESSURE CHRONIC.DISEASE
## 1      0  69      1          2          2          1          1
## 2      0  74      2          1          1          1          2
## 3      1  59      1          1          1          2          1
## 4      0  63      2          2          2          1          1
## 5      1  63      1          2          1          1          1
## 6      1  75      1          2          1          1          2
##      FATIGUE ALLERGY WHEEZING ALCOHOL.CONSUMING COUGHING SHORTNESS.OF.BREATH
## 1      2      1      2          2      2          2
## 2      2      2      1          1      1          2
## 3      2      1      2          1      2          2
## 4      1      1      1          2      1          1
## 5      1      1      2          1      2          2
## 6      2      2      2          1      2          2
##      SWALLOWING.DIFFICULTY CHEST.PAIN LUNG_CANCER
## 1          2          2          1
## 2          2          2          1
## 3          1          2          0
## 4          2          2          0
## 5          1          1          0
## 6          1          1          1

```

- GENDER and LUNG_CANCER were encoded to 0 and 1.

```

# Calculate the totals
total_healthy <- sum(data$LUNG_CANCER == 0)
total_lung_cancer <- sum(data$LUNG_CANCER == 1)

# Create a dataframe
totals_df <- data.frame(
  Category = c("Healthy Patients", "Lung Cancer Patients"),
  Total = c(total_healthy, total_lung_cancer)
)

print(totals_df)

```

```

##              Category Total
## 1    Healthy Patients    39
## 2 Lung Cancer Patients   270

```

```

#Split data into test and train
set.seed(123)


train_index <- sample(1:nrow(data), 0.8 * nrow(data))
train_data <- data[train_index, ]
test_data <- data[-train_index, ]
val_data <- data[-train_index, ] #validation set

```

```

# Compute the correlation matrix
corr_df <- train_data %>% select(-LUNG_CANCER)
cor_matrix <- cor(corr_df)
# Plot correlation matrix
corrplot(cor_matrix, method = "square", tl.cex = 0.6)


```



```

# Ridge solution path
x_tr <- as.matrix(train_data[, -16])
y_tr <- train_data[, 16, drop = T]
x_te <- as.matrix(test_data[, -16])
y_te <- test_data[, 16, drop = T]
std_fit <- preProcess(x_tr, method = c("center", "scale"))
x_tr_std <- predict(std_fit, x_tr)
x_te_std <- predict(std_fit, x_te)
fit_ridge <- glmnet(x_tr_std, y_tr, alpha = 0)
plot_glmnet(fit_ridge)

```




```

# Ridge solution path
#Standardization
set.seed(0)
cv_fit_ridge <- cv.glmnet(x_tr_std, y_tr, alpha = 0, nfolds=5)
r_tr_pred <- predict(cv_fit_ridge, newx = x_tr_std)
r_te_pred <- predict(cv_fit_ridge, newx = x_te_std)
r_tr_error <- mean((r_tr_pred - y_tr)^2)
r_te_error <- mean((r_te_pred - y_te)^2)
best_coefR <- as.vector(coef(cv_fit_ridge))
best_coefR <- c(coef(cv_fit_ridge)[1], best_coefR)
best_coefR <- best_coefR[-1]
variables <- c("GENDER", "SMOKING", "AGE", "YELLOW_FINGERS", "ANXIETY", "PEER_PRESSURE", "CHRONIC.DISEASE", "PEER_PRE", "YELLOW_F")
coef_names <- c("(Intercept)", variables)
ridge_df <- data.frame(coef_names, best_coefR, Method = "Ridge Regression", Test_Error = r_te_error) %>%

```

```
# Lasso Solution Path
fit_lasso <- glmnet(x_tr_std, y_tr)
plot_glmnet(fit_lasso)
```



```
# Lasso Solution Path
cv_fit_lasso <- cv.glmnet(x_tr, y_tr, alpha = 1, nfolds = 5)
l_tr_pred <- predict(cv_fit_lasso, newx = x_tr)
l_te_pred <- predict(cv_fit_lasso, newx = x_te)
l_tr_error <- mean((l_tr_pred - y_tr)^2)
l_te_error <- mean((l_te_pred - y_te)^2)
best_coefL <- as.vector(coef(cv_fit_lasso))
variables <- c("GENDER", "SMOKING", "YELLOW_FINGERS", "ANXIETY", "PEER_PRESSURE", "CHRONIC.DISEASE", "F")
coef_names <- c("(Intercept)", variables)
lasso_df <- data.frame(coef_names, best_coefL, Method = "Lasso", Test_Error = l_te_error) %>% pivot_wider()

rl_errors_df <- data.frame(
  Model = c("Ridge", "Lasso"),
  Testing_Error = c(r_te_error, l_te_error)
)
print(rl_errors_df)
```

```
## Model Testing_Error
## 1 Ridge 0.09986412
## 2 Lasso 0.10405987
```

```

#KNN
# Standardizing variables for optimization
fit_std <- preProcess(train_data, method = "scale")
train_std <- predict(fit_std, newdata = train_data)
test_std <- predict(fit_std, newdata = test_data)

# Sequence of k values for KNN algorithm
k_seq <- seq(from = 1, to = 50, by = 5)
train_error_seq_std <- test_error_seq_std <- NULL

# Fitting KNN model on std predictor variables
for (k in k_seq) {
  fit_knn_std <- knnreg(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE,
    data = train_std, k = k)

  pred_knn_train <- predict(fit_knn_std, newdata = train_std, type = "class")
  train_error_seq_std <- c(train_error_seq_std, mean(pred_knn_train != train_data$LUNG_CANCER))

  pred_knn_test <- predict(fit_knn_std, newdata = test_std, type = "class")
  test_error_seq_std <- c(test_error_seq_std, mean(pred_knn_test != test_data$LUNG_CANCER))
}

# Fitting KNN on unstd predictor variables
train_error_seq <- test_error_seq <- NULL

for (k in k_seq) {
  fit_knn <- knnreg(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE,
    data = train_data, k = k)

  pred_knn_train <- predict(fit_knn, newdata = train_data, type = "class")
  train_error_seq <- c(train_error_seq, mean(pred_knn_train != train_data$LUNG_CANCER))

  pred_knn_test <- predict(fit_knn, newdata = test_data, type = "class")
  test_error_seq <- c(test_error_seq, mean(pred_knn_test != test_data$LUNG_CANCER))
}


# Combining results
knn_std <- data.frame(K = k_seq, error = c(train_error_seq_std, test_error_seq_std),
  type = rep(c("train_std", "test_std"), each = length(k_seq)))

knn_unstd <- data.frame(K = k_seq, error = c(train_error_seq, test_error_seq),
  type = rep(c("train", "test"), each = length(k_seq)))

# Combine standardized and unstandardized results
knn_combined <- rbind(knn_std, knn_unstd)

ggplot(knn_combined, aes(x = K, y = error, color = type)) +
  geom_point(size = 2) +
  geom_line(size = 2)

```

- Based on errors, $k = 7$ will work best

```

#Logistic Regression and KNN
set.seed(0)
# Fit Logistic Regression
fit_lr <- glm(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE + ALLE
fit_lr_pred <- predict(fit_lr, type = "response")

# Calculate errors for Logistic Regression
lr_train_pred <- predict(fit_lr, newdata = train_data, type = "response")
lr_test_pred <- predict(fit_lr, newdata = test_data, type = "response")

train_pred_binary <- ifelse(lr_train_pred >= 0.5, 1, 0)
test_pred_binary <- ifelse(lr_test_pred >= 0.5, 1, 0)

lr_test_error <- mean(train_pred_binary != test_data$LUNG_CANCER)
lr_train_error <- mean(test_pred_binary != train_data$LUNG_CANCER)

# Fit KNN model
fit_knn <- knn3(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE + ALLE

# Calculate errors for KNN
knn_train_pred <- predict(fit_knn, newdata = train_data)
knn_train_error <- mean(knn_train_pred != train_data$LUNG_CANCER)
knn_test_pred <- predict(fit_knn, newdata = test_data)
knn_test_error <- mean(knn_test_pred != test_data$LUNG_CANCER)

# Create dataframe for model errors
model_errors_df <- data.frame(
  Model = c("Logistic Regression", "KNN"),
  Training_Error = c(lr_train_error, knn_train_error),
  Testing_Error = c(lr_test_error, knn_test_error)
)

print(model_errors_df)

```

```

##           Model Training_Error Testing_Error
## 1 Logistic Regression      0.2226721      0.2429150
## 2           KNN          0.6923077      0.7177419

```

```

#Logistic Regression and KNN
#Cross Validation
# Fit and LOOCV for logistic model
mod_formula_seq <- c("LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE")

loocv_lr <- sapply(mod_formula_seq, function(form){
  mod <- as.formula(form)
  fit <- glm(mod, data = train_data)
  cv.glm(train_data, fit)$delta[1]
})

# LOOCV for KNN Model with K=7
mod_formula_seq_knn <- "LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE"
loocv_knn<-sapply(mod_formula_seq_knn, function(form){
  mod <- as.formula(form)
  mean(sapply(1:nrow(train_data), function(j){
    fit <- knn3(mod, data = train_data[-j, ],k=7)
    pred_lung <- predict(fit, newdata = train_data[j, ])
    (train_data$LUNG_CANCER[j] - pred_lung)^2
  })))
})

# Extract LOOCV errors
lr_loocv_error <- loocv_lr[[1]]
knn_loocv_error <- loocv_knn[[1]]

# Create dataframe for model errors
loocv_df <- data.frame(
  Model = c("Logistic Regression", "KNN"),
  LOOCV_Error = c(lr_loocv_error, knn_loocv_error)
)

print(loocv_df)

```

```

##                Model LOOCV_Error
## 1 Logistic Regression 0.07174301
## 2                KNN 0.44591164

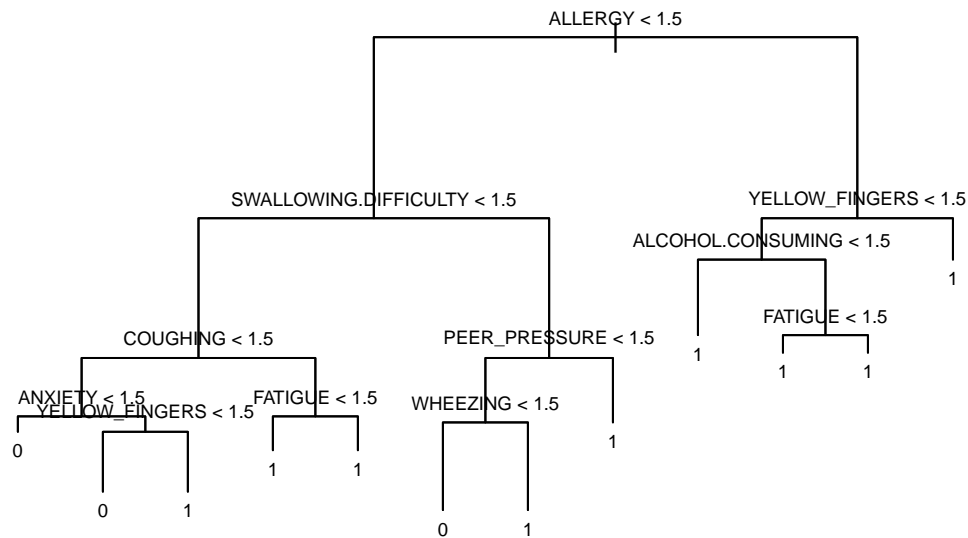
```

- Compared to KNN, logistic has a relatively low misclassification rate.

```

#Decision Tree
# Convert LUNG_CANCER to a factor variable
train_data$LUNG_CANCER <- as.factor(train_data$LUNG_CANCER)
lung_cancer_tree <- tree(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATI
                        data = train_data)
plot(lung_cancer_tree)
text(lung_cancer_tree, cex= 0.6)

```



```

#Decision Tree
# Convert LUNG_CANCER to a numeric variable for error calculation
train_data$LUNG_CANCER <- as.numeric(as.character(train_data$LUNG_CANCER))
test_data$LUNG_CANCER <- as.numeric(as.character(test_data$LUNG_CANCER))
train_predictions <- predict(lung_cancer_tree, newdata = train_data)
test_predictions <- predict(lung_cancer_tree, newdata = test_data)
# Calculate errors
dt_train_error <- mean((train_predictions - train_data$LUNG_CANCER)^2)
dt_test_error <- mean((test_predictions - test_data$LUNG_CANCER)^2)
#Dataframe of errors
dt_errors_df <- data.frame(Model = c("Decision Tree"),
    Training_Error = c(dt_train_error), Testing_Error = c(dt_test_error)
)
print(dt_errors_df)

```

```

##           Model Training_Error Testing_Error
## 1 Decision Tree      0.4541269      0.4560601


```

```

#Decision Tree
# Cross Validation
set.seed(0)
cv_lung_cancer <- cv.tree(lung_cancer_tree)
cv_lung_cancer_df <- data.frame(size = cv_lung_cancer$size, deviance = cv_lung_cancer$dev)
best_size <- cv_lung_cancer$size[which.min(cv_lung_cancer$dev)]

ggplot(cv_lung_cancer_df, aes(x = size, y = deviance)) +
  geom_point(size = 3) +
  geom_line() +
  geom_vline(xintercept = best_size, col = "red")

```



```

cat('CV leads to the optimal tree size as ', best_size, '\n')

```

```

## CV leads to the optimal tree size as 12

```

```

# Random Forest
# Fit random forest model
set.seed(1)
rf_mod_type <- randomForest(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + 
                             data = train_data, ntreeTry = 1000, importance = TRUE)

rf_train_pred_type <- predict(rf_mod_type, newdata = train_data, type = "response")
rf_test_pred_type <- predict(rf_mod_type, newdata = test_data, type = "response")

# Convert predicted values to factor (binary classification)
rf_train_pred_class <- as.factor(ifelse(rf_train_pred_type > 0.5, "1", "0"))
rf_test_pred_class <- as.factor(ifelse(rf_test_pred_type > 0.5, "1", "0"))

# Calculate errors
rf_train_error <- mean(rf_train_pred_class != train_data$LUNG_CANCER)
rf_test_error <- mean(rf_test_pred_class != test_data$LUNG_CANCER)

# Dataframe for errors
rf_errors_df <- data.frame(
  Model = c("Random Forrest"),
  Training_Error = c(rf_train_error),
  Testing_Error = c(rf_test_error)
)

rf_errors_df

```

```

##           Model Training_Error Testing_Error
## 1 Random Forrest      0.04048583      0.1129032

```

```

# Random Forest
# Feature importance for random forest model
rf_importance <- importance(rf_mod_type)
rf_fn <- rownames(rf_importance)
rf_data <- data.frame(Feature = rf_fn, Importance = rf_importance[, 1])
rf_plot <- ggplot(rf_data, aes(x = Importance, y = reorder(Feature, Importance))) +
  geom_bar(stat = "identity", fill = "lightgreen") +
  labs(title = NULL, x = NULL, y = NULL) +
  theme_minimal()

importance(rf_mod_type)


```

```

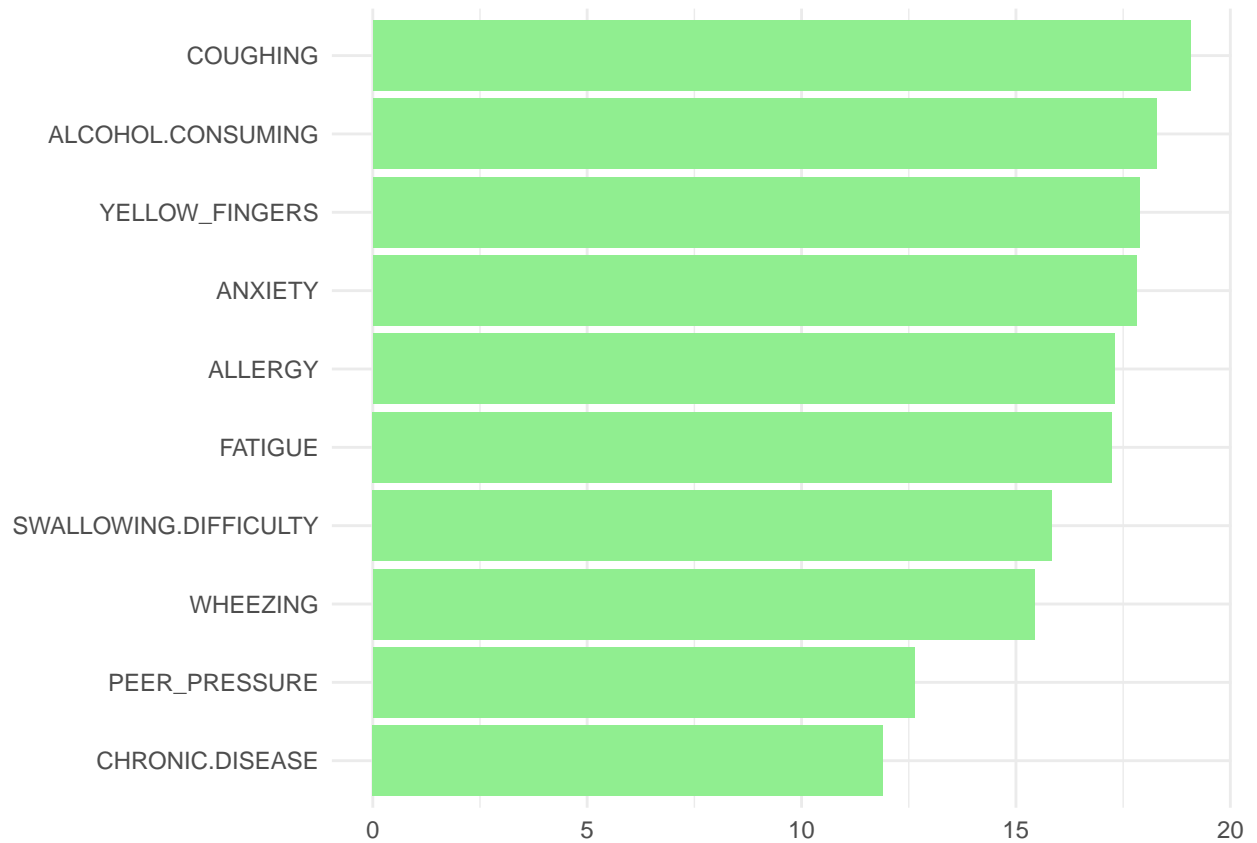
##           %IncMSE IncNodePurity
## YELLOW_FINGERS      17.87960      1.654784
## ANXIETY              17.80993      1.565168
## PEER_PRESSURE       12.64067      1.362142
## CHRONIC.DISEASE     11.89969      1.456343
## FATIGUE              17.24007      1.593161
## ALLERGY              17.29679      2.273151
## WHEEZING             15.43582      1.629250
## ALCOHOL.CONSUMING    18.28042      1.772893
## COUGHING             19.07052      1.957141
## SWALLOWING.DIFFICULTY 15.83987      1.525005

```

```
varImpPlot(rf_mod_type)
```



```
print(rf_plot)
```

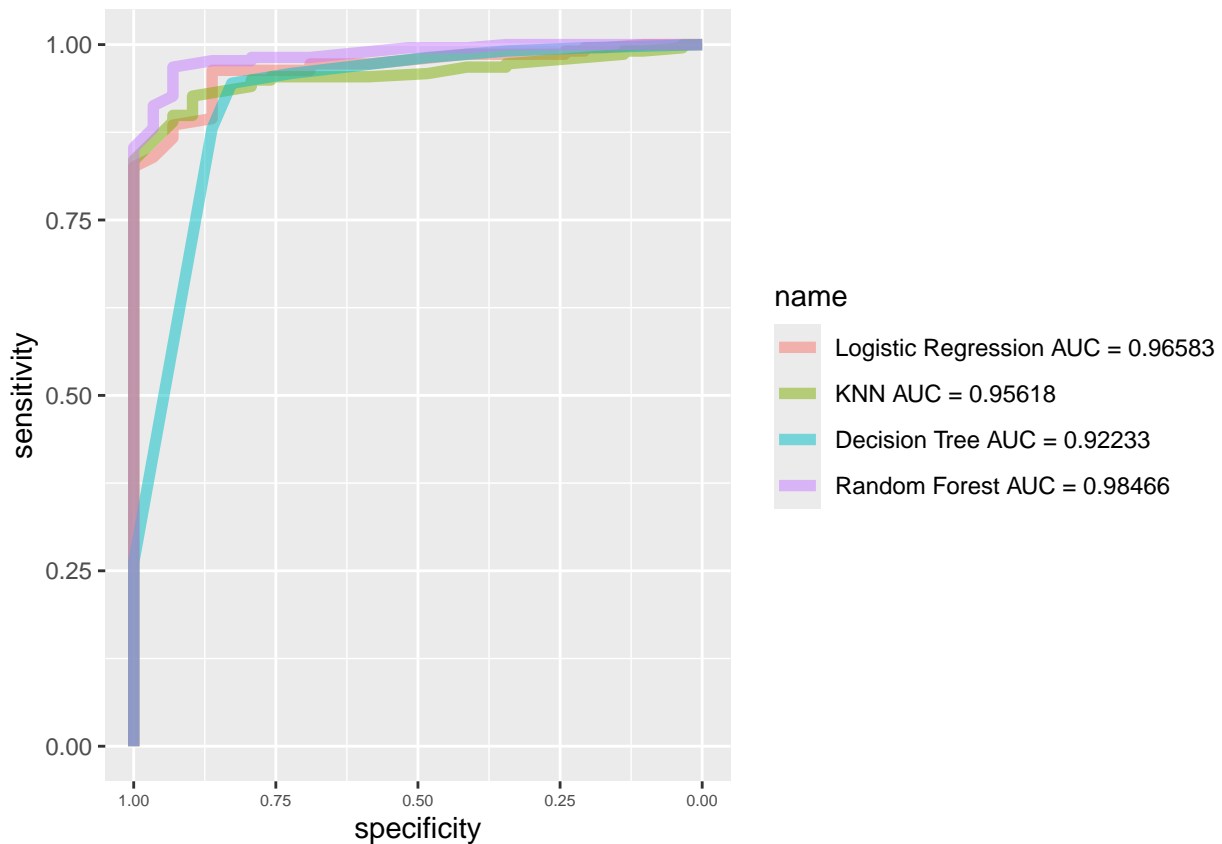


- The most important features across random forest were COUGHING, ALCOHOL CONSUMING, and YELLOW FINGERS.


```

#ROC/AUC PLOT
# Fit Logistic Regression
fit_lr <- glm(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE + ALLE
fit_lr_pred <- predict(fit_lr, type = "response")
roc_lr <- roc(train_data$LUNG_CANCER, fit_lr_pred)
# Fit KNN model
fit_knn <- knn3(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE + AL
fit_knn_pred <- predict(fit_knn, newdata = train_data, type = "prob")
roc_knn <- roc(train_data$LUNG_CANCER, fit_knn_pred[, 2])
# Fit Decision Tree model
fit_tree <- tree(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE + AL
fit_tree_pred <- predict(fit_tree, newdata = train_data, type = "vector")
roc_tree <- roc(train_data$LUNG_CANCER, fit_tree_pred)
# Fit Random Forest model
fit_rf <- randomForest(LUNG_CANCER ~ YELLOW_FINGERS + ANXIETY + PEER_PRESSURE + CHRONIC.DISEASE + FATIGUE
fit_rf_pred <- predict(fit_rf, newdata = train_data, type = "response")
roc_rf <- roc(train_data$LUNG_CANCER, fit_rf_pred)
# Plot ROC curves
rocobjs <- list("Logistic Regression" = roc_lr, "KNN" = roc_knn, "Decision Tree" = roc_tree, "Random Fo
methods_auc <- paste(c("Logistic Regression", "KNN", "Decision Tree", "Random Forest"), "AUC =", round(
ggroc(rocobjs, size = 2, alpha = 0.5) +
  scale_color_discrete(labels = methods_auc) +
  theme(axis.text.x = element_text(size = 6))

```



- Random forest had the highest AUC value at 0.984 followed by logistic regression at 0.97.

```

#Confusion matrix for logistic regression and random forrest
#Logistic Regression
predicted_classes <- ifelse(lr_test_pred >= 0.5, "1", "0")
conf_matrix_lr <- confusionMatrix(as.factor(predicted_classes), as.factor(test_data$LUNG_CANCER), posit
print("Confusion Matrix for Logistic Regression:")

```

```
## [1] "Confusion Matrix for Logistic Regression:"
```

```
print(conf_matrix_lr)
```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0   1
##           0   6   1
##           1   4  51
##
##              Accuracy : 0.9194
##              95% CI : (0.8217, 0.9733)
##      No Information Rate : 0.8387
##      P-Value [Acc > NIR] : 0.05172
##
##              Kappa : 0.6608
##
##  McNemar's Test P-Value : 0.37109
##
##              Sensitivity : 0.9808
##              Specificity : 0.6000
##              Pos Pred Value : 0.9273
##              Neg Pred Value : 0.8571
##              Prevalence : 0.8387
##              Detection Rate : 0.8226
##      Detection Prevalence : 0.8871
##              Balanced Accuracy : 0.7904
##
##              'Positive' Class : 1
##

```

```

#Random Forrest
predicted_classes2 <- ifelse(rf_test_pred_type >= 0.5, "1", "0")
conf_matrix_rf <- confusionMatrix(as.factor(predicted_classes2), as.factor(test_data$LUNG_CANCER), posi
print("Confusion Matrix for Random Forrest:")

```

```
## [1] "Confusion Matrix for Random Forrest:"
```

```
print(conf_matrix_rf)
```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0   1

```

```

##          0  4  1
##          1  6 51
##
##          Accuracy : 0.8871
##          95% CI : (0.7811, 0.9534)
##    No Information Rate : 0.8387
##    P-Value [Acc > NIR] : 0.1967
##
##          Kappa : 0.4771
##
##    McNemar's Test P-Value : 0.1306
##
##          Sensitivity : 0.9808
##          Specificity : 0.4000
##    Pos Pred Value : 0.8947
##    Neg Pred Value : 0.8000
##          Prevalence : 0.8387
##    Detection Rate : 0.8226
##    Detection Prevalence : 0.9194
##    Balanced Accuracy : 0.6904
##
##    'Positive' Class : 1
##

```

- Both logistic regression and random forest models had high accuracy rates at 92% and 89%.